

Tomas Sandnes
tomas.sandnes@kristiania.no

Week 5

DB1102 / PGR 111 – DATABASES



Today's topics

(Today's chapters: 5.1-5.3 in Norwegian book, 3.8 & 4.2 in English book)

- Exercise walkthrough of task 5 from last week
- Subqueries
- VIEWS



Exercise walkthrough

Exercise walkthrough, 1 task from week 4

- We're looking into task 5 from the previous week:
 - *Retrieve all the countries in the world, and all their cities, where the population of the country is less than 1000 people. We want to retrieve the name of the country (AS CountryName), the population of the country (AS CountryPop), which continent the country belongs to, the name of all cities (AS CityName) and the population of the cities (AS CityPop). We want the result sorted alphabetically by the country names. (Hint: My answer returns 10 rows, is that the same number of rows you get?)*

Subqueries

Subqueries

- As mentioned in previous lessons, the result of a `SELECT` is presented as an SQL table.
 - It forms columns and rows the same way as existing tables in the database do.
- Therefore, it is ok to use the result of one `SELECT` as part of another SQL statement. Using **a `SELECT` statement inside another SQL statement** in this way is called a **subquery**.
 - (For the exercises on table `mail` and `person` in a previous week, you used a `SELECT` as part of an `INSERT INTO` statement.)

Subqueries – cont.

- Sometimes we need the answer from one query to complete another SQL statement.
 - We have already done this regarding insert into / update / delete from.
 - We can also do this when only SELECT statements are involved.
- *Example:* We want to figure out how many cities that have a population equal to or greater than the average population.
 - Before we can complete this query, we need to figure out what the average city population is.

Subqueries – cont.

- Task: *"How many cities have a population equal to or greater than the average city population?"*
 - We can solve it with a subquery (a SELECT) inside a SELECT, like this:

```
SELECT COUNT(*)  
FROM city  
WHERE Population >=  
(SELECT AVG(Population) FROM city);
```


Subqueries – cont.

- Another one: *"List all European cities"*
 - We can solve it like this:

```
SELECT Name
FROM city
WHERE CountryCode IN
(SELECT Code
 FROM country
 WHERE Continent = 'Europe');
```

- Explanation of IN: Lets us compare a column with a group of values.
- (Note: This one could have been solved with a JOIN as well.)

Subqueries – cont.

- Yet another example: *"What size, as a percentage of the world's largest country, does each country have?"*
 - We can solve it with a subquery (a SELECT) inside a SELECT, like this:

```
SELECT Name, SurfaceArea / (SELECT MAX(SurfaceArea) FROM country) * 100
FROM country
ORDER BY SurfaceArea DESC;
```


- Or what about: *"We want to see what percentage the inhabitants of a country make up of the total population of Earth."*
 - This is one of today's exercises. :-)

VIEW

What is view?

- A **VIEW** is a pre-created query for one or more tables.
 - *Example:* We only want ID, NAME, and ADDRESS, for "STAVANGER"

ID	NAME	ADDRESS	CITY	POSITION	SALARY	HIREDATE	BNO
3	JON	BRUVEIEN 7	STAVANGER	MANAGER	70000	08-SEP-15	1
2	MARIE	STRILEGATEN 8	BERGEN	MANAGER	60000	01-JAN-15	2
1	SUSANNE	SKRAPLODDVEIEN 62	OSLO	MANAGER	90000	01-JUL-14	3
20	OLAV	GALMANNNSVEIEN 4	STAVANGER	BROKER	52000	07-JUL-17	1
5	DAVID	GULERLEVEIEN 43	STAVANGER	SECRETARY	36000	14-JUN-15	1
4	ANNE	STRANDGATEN 5	STAVANGER	BROKER	24000	12-DEC-16	1
8	GUSTAV	NORDLYSVEIEN 78	STAVANGER	BROKER			
9	OLAVA	LOMVIVEIEN 57	STAVANGER	BROKER			
11	LEONORA	RØDVEIEN 6	OSLO	BROKER			
10	TEODOR	TULIPAN 12	OSLO	SECRETARY			
6	JONNAS	KIRKEVEIEN 7	BERGEN	BROKER			
12	TULLA	BLÅVEIEN 7a	OSLO	BROKER			
18	FREDRIK	LASSOVEIEN 37	OSLO	BROKER			
7	KARL	OLAVSGATE 7	OSLO	BROKER			
16	SMUKKA	GRAUTSTIEN 43	OSLO	BROKER			
17	KARL	BLÅVEISVEIEN 7	OSLO	BROKER	41000	12-MAY-96	3



ID	NAME	ADDRESS
3	JON	BRUVEIEN 7
20	OLAV	GALMANNNSVEIEN 4
5	DAVID	GULERLEVEIEN 43
4	ANNE	STRANDGATEN 5
8	GUSTAV	NORDLYSVEIEN 78
9	OLAVA	LOMVIVEIEN 57

Why use views?

- Security: Restrict data access in the database.
- Reduce complexity when writing SQL queries later:
 - Can make complex queries easier.
 - Reduces the amount of data for the user.
- Customization: Achieve different views of the database.
 - user groups / applications

Disadvantages of views

- More complexity related to updates and changes:
 - There are restrictions on when you can change underlying data through a view.
 - The structure is determined when the VIEW is created, and will not automatically change later (VIEW based on `SELECT * FROM ...`)
- Performance:
 - A view can join many tables, and thus be a relatively heavy query.
 - This is not always easy to see for the user.

Create a view

```
CREATE [OR REPLACE] VIEW view_name AS  
SELECT subquery;
```

- OR REPLACE → overwrites a VIEW if it already exists
- subquery → the SELECT query we want to store

Example, the world schema

```
CREATE OR REPLACE VIEW EuropeCountry_view AS  
SELECT Code, Name, Population  
FROM country  
WHERE Continent = 'Europe';
```

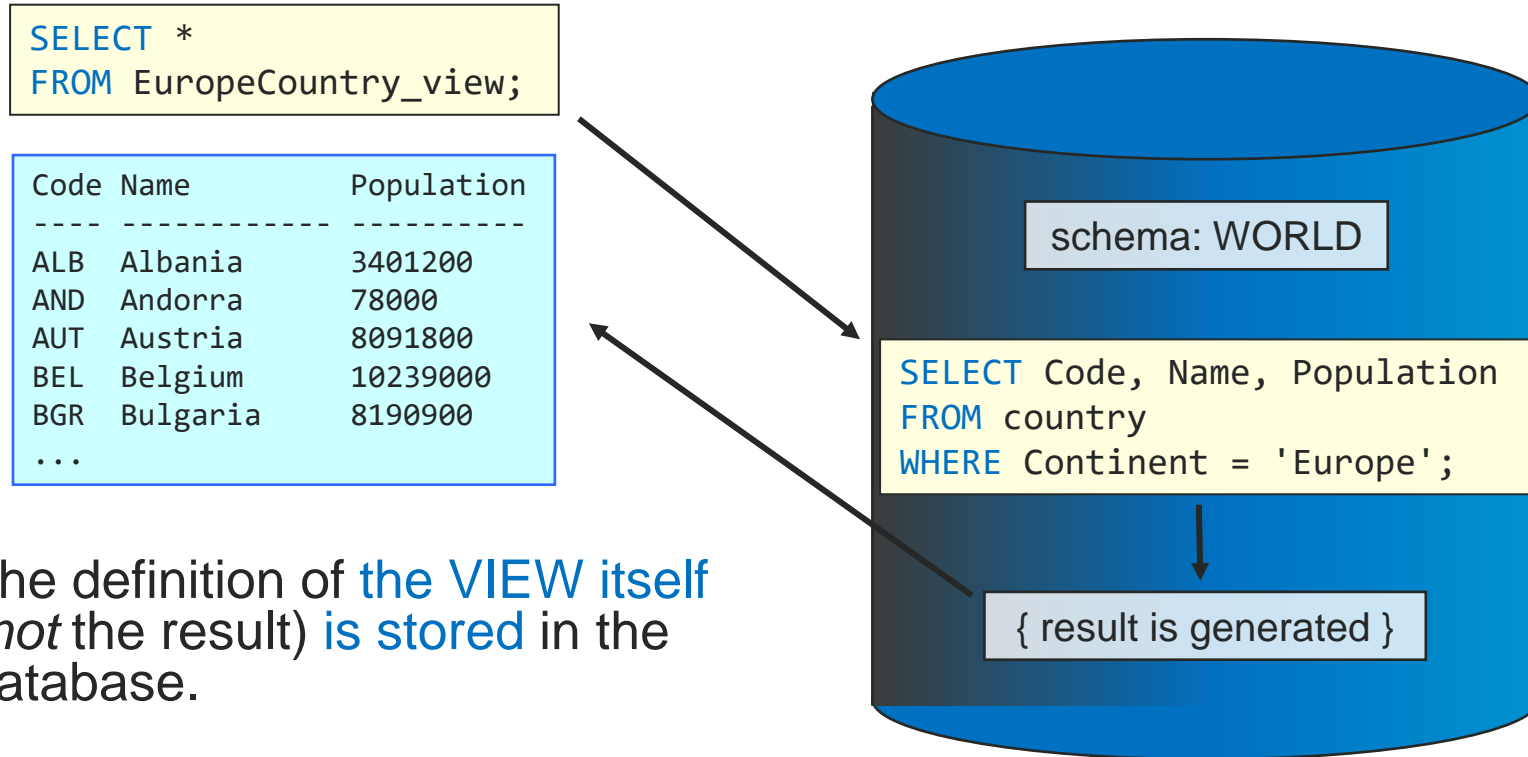
```
SELECT *  
FROM EuropeCountry_view;
```



Code	Name	Population
ALB	Albania	3401200
AND	Andorra	78000
AUT	Austria	8091800
BEL	Belgium	10239000
BGR	Bulgaria	8190900
BIH	Bosnia and Herzegovina	3972000
BLR	Belarus	10236000
CHE	Switzerland	7160400
CZE	Czech Republic	10278100
DEU	Germany	82164700
DNK	Denmark	5330000
ESP	Spain	39441700
...		

- Views are treated the same way as tables are.

This is how querying a view works



- The definition of **the VIEW itself** (not the result) **is stored** in the database.

Example, some variations

- We can give views their own column names:

```
CREATE OR REPLACE VIEW EuropeCountry_view AS  
SELECT Code AS ID, Name AS Country, Population  
FROM country  
WHERE Continent = 'Europe';
```



```
CREATE OR REPLACE VIEW EuropeCountry_view  
(ID, Country, Population) AS  
SELECT Code, Name, Population  
FROM country  
WHERE Continent = 'Europe';
```

```
SELECT *  
FROM EuropeCountry_view;
```



ID	Country	Population
----	-----	-----
ALB	Albania	3401200
AND	Andorra	78000
AUT	Austria	8091800
BEL	Belgium	10239000
BGR	Bulgaria	8190900
BIH	Bosnia and Herzegovina	3972000
BLR	Belarus	10236000
CHE	Switzerland	7160400
CZE	Czech Republic	10278100
DEU	Germany	82164700
DNK	Denmark	5330000
ESP	Spain	39441700
...		

Group functions used in views

```
CREATE OR REPLACE VIEW ContinentPopulation_view AS  
SELECT Continent, SUM(Population) AS Population  
FROM country  
GROUP BY Continent  
ORDER BY SUM(Population) ASC;
```

```
SELECT *  
FROM ContinentPopulation_view;
```



Continent	Population
-----	-----
Antarctica	0
Oceania	30401150
South America	345780000
North America	482993000
Europe	730074600
Africa	784475000
Asia	3705025700

- Group functions can be used in views.
 - It is a good idea to use AS (alias) to customize the column names.

Exercise

- Create a view to display the cities in the world with population > 5 million.
 - It should display city, population and which country the city is located in.
 - It should be sorted on population descending.

```
CREATE OR REPLACE VIEW LargeCities_view
(city, population, country) AS
SELECT city.name, city.population, country.name
FROM city
JOIN country ON CountryCode = Code
WHERE city.Population > 5000000
ORDER BY city.Population DESC;
```

	city	population	country
▶	Mumbai (Bombay)	10500000	India
	Seoul	9981619	South Korea
	São Paulo	9968485	Brazil
	Shanghai	9696300	China
	Jakarta	9604900	Indonesia
	Karachi	9269265	Pakistan
	Istanbul	8787958	Turkey

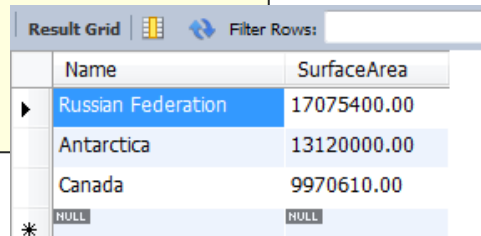
- Then, lets retrieve the 7 largest cities:

```
SELECT *
FROM largeCities_view
LIMIT 7;
```

New SQL word: LIMIT

- **LIMIT** limits (d'oh!) the number of rows in your result.
- Example:
 - I want to find the three physically largest countries in the world.

```
SELECT Name, SurfaceArea  
FROM Country  
ORDER BY SurfaceArea DESC  
LIMIT 3;
```



	Name	SurfaceArea
▶	Russian Federation	17075400.00
	Antarctica	13120000.00
	Canada	9970610.00
	NULL	NULL

Tips for using CREATE VIEW

- Same **tip for creating VIEW** as for doing insert / update / delete based on queries:
 - Always **test the SELECT first**, so you know if it works before you use it in the view!

```
CREATE OR REPLACE VIEW largeCities
(city, population, country)
AS SELECT ci.name, ci.population, co.name
FROM city ci LEFT JOIN country co ON ci.countryCode = co.Code
WHERE ci.Population>5000000
ORDER BY population DESC;
```

Updates through views

- Views can be used to update data in an underlying table.
- *Note:* There are ISO **restrictions** that must be followed **before a view can be used for updates**:
 - The view can only reference **one table**.
 - **DISTINCT cannot** be part of the view.
 - **All elements** in the select part of the view **must be columns** (not constants, summations, etc.)
 - **No GROUP BY** or **HAVING**.
 - Rows that are added **must follow the integrity rules** for the underlying table (not null, etc.).

Updates through views – cont.

- Note that a view updates the data in the table! (Not just the view itself.)

```
UPDATE EuropeCountry_view  
SET ID = 'A_Z'  
WHERE Country = 'Austria';
```

```
SELECT ID, Country  
FROM EuropeCountry_view;
```



ID	Country
----	-----
ALB	Albania
AND	Andorra
A_Z	Austria
BEL	Belgium
BGR	Bulgaria
...	

```
SELECT Code, Name, Population  
FROM country  
WHERE Name = 'Austria';
```



CODE	Name	Population
----	-----	-----
A_Z	Austria	8091800

Deleting views

- Syntax to delete a view is almost the same as for a table:
 - Just replace the word "table" with "view".

```
DROP VIEW EuropeCountry_view;
```

Today's exercises & next lesson

- Now: 2 hours of exercises.
- Exercises are found on Canvas. Short summary:
 - Try to make some subqueries on your own.
 - Combine what we just learned about VIEW with SQL knowledge syntax from our earlier lessons.
- Main contents for the next lesson:
 - Not new content, but a sum-up of what we have learned so far!
 - Walkthrough of exercises *YOU* want to see solved from previous weeks. Send me those week- & task-numbers. :-)

The

